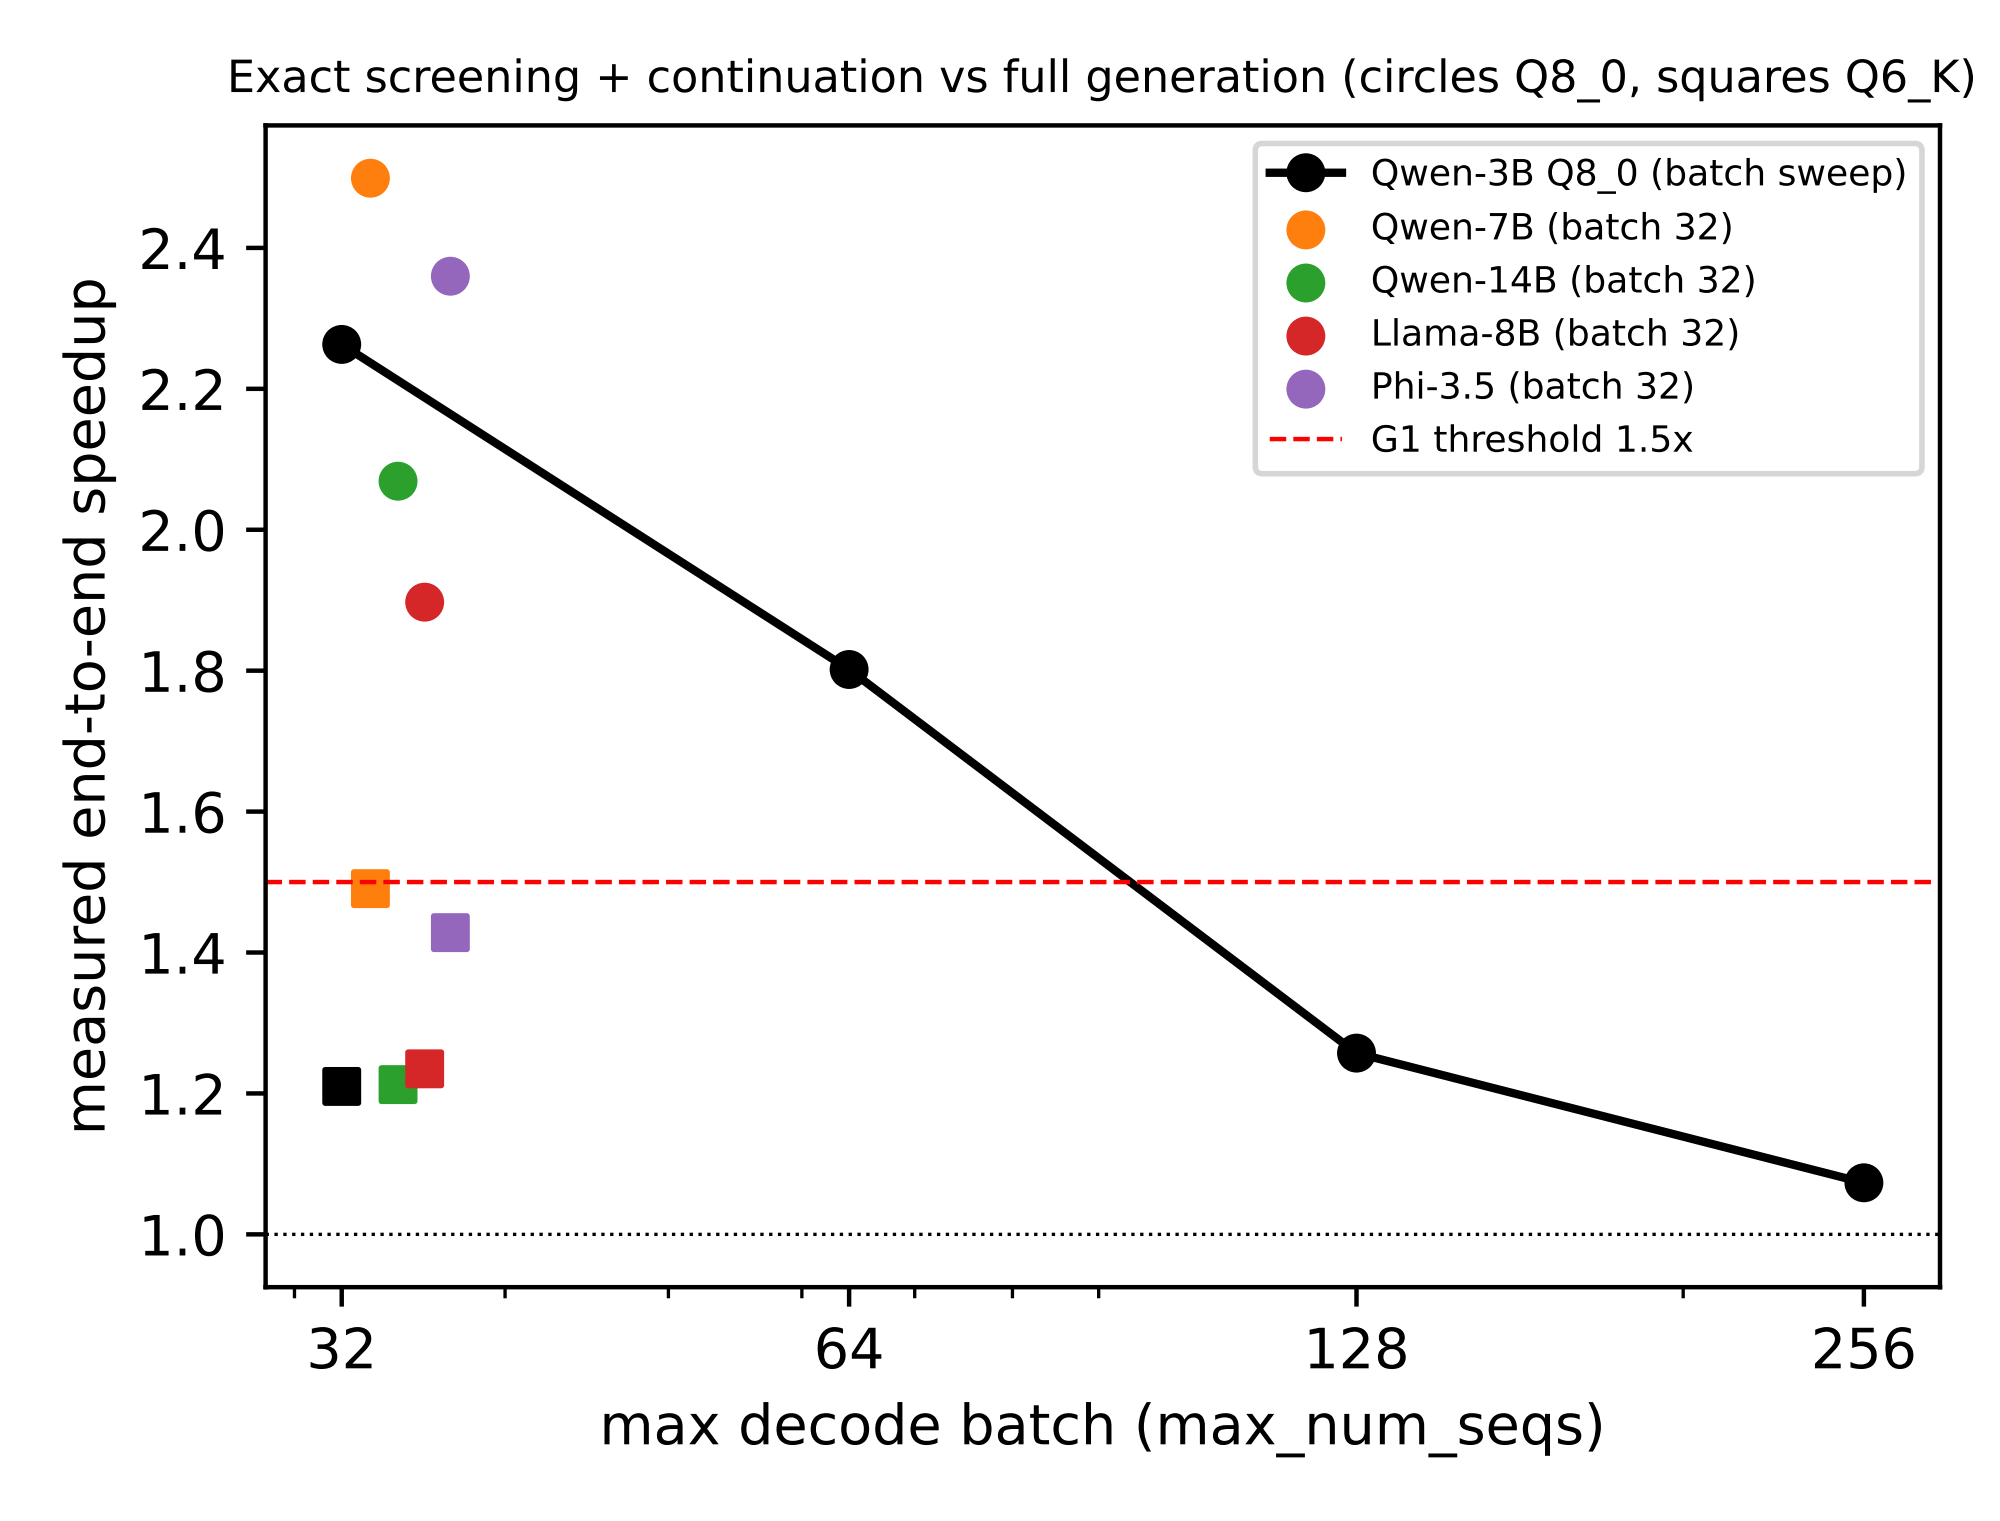
Exact screening + continuation vs full generation (circles Q8_0, squares Q6_K)
measured end-to-end speedup
max decode batch (max_num_seqs)
2.4
2.2
2.0
1.8
1.6
1.4
1.2
1.0
32
64
128
256
Qwen-3B Q8_0 (batch sweep)
Qwen-7B (batch 32)
Qwen-14B (batch 32)
Llama-8B (batch 32)
Phi-3.5 (batch 32)
G1 threshold 1.5x